

Machine learning

Aheer srabon

April 1, 2024

1 Introduction

2 Linear regression with one variable

2.1 Model and cost function

2.2 Parameter learning

3 Linear regression with multiple variables

3.1 Multivariate linear regression

3.2 Computing parameters analytically

4 Logistic regression

4.1 Classification and representation

4.2 Logistic regression model

4.3 Multiclass classification

5 Regularization

Regularization helps prevent models from overfitting the training data.

- 5.1 Solving the problem of overfitting
- 6 Neural networks: representation
- 7 Neural networks: learning
 - 7.1 Cost function and backpropagation
 - 7.2 Backpropagation in practice
 - 7.3 Application of neural networks
- 8 Advice for applying machine learning
 - 8.1 Evaluating a learning algorithm
 - 8.2 Bias vs. variance
- 9 Machine learning system design
 - 9.1 Building a spam classifier
 - 9.2 Handling skewed data
 - 9.3 Using large data sets
- 10 Support vector machines
 - 10.1 Large margin classification
 - 10.2 Kernels
 - 10.3 SVMs in practice
- 11 Unsupervised learning

While supervised learning algorithms need labeled examples (x, y) , unsupervised learning algorithms need only the input (x) . Clustering is used for market segmentation, text summarization, among many other applications.

Principal Component Analysis (PCA) is used to speed up learning algorithms, and is sometimes incredibly useful for visualizing and helping to understand the data.

- 11.1 Clustering
- 12 Dimensionality reduction
 - 12.1 Principal Component Analysis
 - 12.2 Applying PCA
- 13 Anomaly detection
 - 13.1 Density estimation
 - 13.2 Building an anomaly detection system
 - 13.3 Multivariate gaussian distribution
- 14 Recommender systems
 - 14.1 Predicting movie ratings
 - 14.2 Collaborative filtering
 - 14.3 Low rank matrix factorization
- 15 Large scale machine learning
 - 15.1 Gradient descent with large datasets
 - 15.2 Advanced topics
- 16 Application example: photo OCR